

Nonlinear least-squares, Gauss-Newton, Levenberg-Marquardt, and connections...

Tuesday, September 23, 2025 2:19 PM

Solve: $f_i(\vec{x}) = 0 \quad \left. \begin{array}{l} \\ \vdots \\ f_m(\vec{x}) = 0 \end{array} \right\} m \text{ equations}$
 i.e. $\vec{F}(\vec{x}) = \vec{0}, \quad F(\vec{x}) = \begin{bmatrix} f_1(\vec{x}) \\ \vdots \\ f_m(\vec{x}) \end{bmatrix}$

$F: \mathbb{R}^n \rightarrow \mathbb{R}^m$

$x \in \mathbb{R}^n$

n variables/
"Unknowns"

So far we took $m=n$... what if we relax that?

$m < n$ often there are multiple solutions

$m > n$ often there's no solution } let's focus on this

→ No solution... next best thing is often the least-squares solution:

* $\min_{x \in \mathbb{R}^n} \left(f(\vec{x}) := \frac{1}{2} \sum_{i=1}^m f_i(\vec{x})^2 \right) \quad f: \mathbb{R}^n \rightarrow \mathbb{R}$

Optimization: $\min_{x \in \mathbb{R}^n} f(\vec{x})$

Canonical methods: ① gradient descent,

$\vec{x}^{(k+1)} = \vec{x}^{(k)} - \eta \cdot \nabla f(\vec{x}^{(k)})$ scalar stepsize

② Newton's method (for minimization)

$\vec{x}^{(k+1)} = \vec{x}^{(k)} - \nabla^2 f(\vec{x}^{(k)})^{-1} \cdot \nabla f(\vec{x}^{(k)})$

Apply these to our least-squares problem J or \bar{J} or $J(\vec{x})$ is the Jacobian of F

① gradient descent,

$\nabla f(\vec{x}) = \frac{1}{2} \sum_{i=1}^m 2 \cdot f_i(\vec{x}) \cdot \nabla f_i(\vec{x}) = \bar{J}^T \cdot F(\vec{x})$

$J = \begin{bmatrix} -\nabla f_1^T \\ \vdots \\ -\nabla f_m^T \end{bmatrix}$

so $\vec{x}^{(k+1)} = \vec{x}^{(k)} - \eta \cdot \bar{J}(\vec{x}^{(k)})^T \cdot F(\vec{x}^{(k)}) = \begin{bmatrix} 1 \\ \nabla f_1 \dots \nabla f_m \\ 1 \end{bmatrix} \cdot \begin{bmatrix} f_1 \\ \vdots \\ f_m \end{bmatrix}$

② Newton's method (for minimization)

$\nabla^2 f(\vec{x}) = \sum_{i=1}^m f_i(\vec{x}) \cdot \nabla^2 f_i(\vec{x}) + \underbrace{\nabla f_i(\vec{x}) \cdot \nabla f_i(\vec{x})^T}_{\text{Scalar}} \in \mathbb{R}^{n \times n}$

so $\vec{x}^{(k+1)} = \vec{x}^{(k)} - \nabla^2 f(\vec{x}^{(k)})^{-1} \cdot \bar{J}(\vec{x}^{(k)})^T \cdot F(\vec{x}^{(k)}) = \bar{J}^T \cdot \bar{J}$

This is equivalent to Newton's method (for root finding) applied to $\nabla f(\vec{x}) = \vec{0}$

Not equivalent to Newton's method (for root finding) applied to $F(\vec{x}) = \vec{0}$

③ (new) "Gauss-Newton", $\vec{x}^{(k+1)} = \vec{x}^{(k)} - \eta \cdot (\bar{J}^T \bar{J})^{-1} \bar{J}^T \cdot F(\vec{x}^{(k)})$

motivation 1: do Newton (for minimization) but approximate $\nabla^2 f(\vec{x})$ with just this term!
 Saves needing to find $\nabla^2 f_i$, and you already needed \bar{J} anyway

motivation 2: $\vec{x}^{(k+1)} = \underset{\vec{x} \in \mathbb{R}^n}{\operatorname{argmin}} \frac{1}{2} \sum_{i=1}^m (f_i(\vec{x}^{(k)}) + \nabla f_i(\vec{x}^{(k)})^T (\vec{x} - \vec{x}^{(k)}))^2$
 linearize inside the square

Nonlinear least-squares, p. 2

Tuesday, September 23, 2025 5:17 PM

(3) Levenberg - Marquardt (not correctly described in our book)

is a robust version of Gauss-Newton, suitable for real problems.

Common in software (just don't confuse with linear least-squares methods)
for nonlinear least-squares

Comparison

let $J = J(\vec{x}) \in \mathbb{R}^{m \times n}$, $F: \mathbb{R}^n \rightarrow \mathbb{R}^m$. Solve $\boxed{F(\vec{x}) = \vec{0}}$

$m \geq n$, nonlinear least squares, define $f(\vec{x}) = \frac{1}{2} \sum_{i=1}^m f_i(\vec{x})^2$

(1) Gradient descent
 $\vec{x} \leftarrow \vec{x} - \eta \cdot \nabla f(\vec{x})$

$m=n$, directly solve $F(\vec{x}) = \vec{0}$

(2) Fixed point iteration

$$\vec{x} \leftarrow \vec{x} - \eta \cdot F(\vec{x})$$

\nwarrow positive stepsize, chosen to (hopefully) make contractive

(2) Newton (for optimization)

$$\vec{x} \leftarrow \vec{x} - \nabla^2 f(\vec{x})^{-1} \cdot \nabla F(\vec{x})$$

(b) Newton's Method (for root-finding) aka Newton-Raphson

(3) Gauss-Newton

$$\vec{x} \leftarrow \vec{x} - (J^T J)^{-1} J^T F(\vec{x})$$

$$\vec{x} \leftarrow \vec{x} - J^{-1} \cdot F(\vec{x})$$

$n \times n$ so inverse makes sense

if $m > n$ and J has rank n \square
then $(J^T J)^{-1} J^T = J^+$

the Moore-Penrose pseudoinverse.
"pinv" in Matlab / numpy.linalg, but
better to use np.linalg.lstsq

if $m=n$ and J invertible,
Gauss-Newton is Newton (root-finding)

$$\text{Since } (J^T J)^{-1} J^T = J^{-1} J^{-T} J^T = J^{-1}$$

...and equivalent to Newton for optimization if F is affine

nonlinear least-squares:

under mild assumptions, a "solution"
always exists, but might not
be a solution to $F(\vec{x}) = \vec{0}$.

root-finding: root may

not exist! i.e.

equations could be incompatible /
inconsistent

Also, have issues of local min vs. global min

Both approaches:

- might need to initialize close
- may have singular or ill-conditioned matrices to invert
- (2), (3), (b) scale $O(n^3)$ w/ dimension n
- (1), (2) often better in high dimensions



Scipy.linalg.lstsq is for linear least squares (e.g. QR factorizations, etc.)

Scipy.optimize.least_squares is for nonlinear least squares
(probably uses Levenberg - Marquardt)

DON'T CONFUSE
THE TWO